

MLSS 2019 London

Kernels Part III: Large scale

Lorenzo Rosasco

MaLGa- Machine learning Genova Center

Universit'a di Genova

MIT

IIT

Outline

Large scale

Random features

Nyström

Optimization

Conclusion

Learning with kernels


$$\hat{\mathbf{f}}_{\lambda}(\mathbf{x}) = \sum_{i=1}^n \mathbf{k}(\mathbf{x}, \mathbf{x}_i) \mathbf{c}_i, \quad \mathbf{c} = (\hat{\mathbf{K}} + \lambda n \mathbf{I})^{-1} \hat{\mathbf{y}}$$

$$(\hat{\mathbf{K}})_{ij} = \mathbf{k}(\mathbf{x}_i, \mathbf{x}_j)$$

Requires¹: time $O(c_k n^2 + n^3)$ space $O(n^2)$.

Bottlenecks for kernels

$$\hat{\mathbf{f}}_{\lambda}(\mathbf{x}) = \sum_{i=1}^n \mathbf{k}(\mathbf{x}, \mathbf{x}_i) \mathbf{c}_i$$
$$(\hat{\mathbf{K}} + \lambda n \mathbf{I}) \mathbf{c} = \hat{\mathbf{y}}$$



Computations: **Space** $O(n^2)$ **Kernel eval.** $O(n^2)$ **Time** $O(n^3)$

Kernel space//time

Computations: **Space** $O(n^2)$ **Kernel eval.** $O(n^2)$ **Time** $O(n^3)$

- ▶ Kernel methods require manipulating \hat{K} .
- ▶ Memory is the main bottleneck.
- ▶ On the fly kernel evaluation helps but does not solve the problem.

Outline

Large scale

Random features

Nystrom

Optimization

Conclusion

Feature maps

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') = \Phi(\mathbf{x})^\top \Phi(\mathbf{x}')$$

$$\hat{K} = \underbrace{\hat{\Phi}}_{p = \infty} \hat{\Phi}^\top$$

$$p \gg n$$

LS with features

Let $\hat{\Phi} \in \mathbb{R}^{np}$, $p \leq \infty$ Then,

$$\hat{\mathbf{w}}_{\lambda} = (\hat{\Phi}^{\top} \hat{\Phi} + \lambda n \mathbf{I})^{-1} \hat{\Phi}^{\top} \hat{\mathbf{y}}$$

Requires: time $O(np^2 + p^3)$ space $O(n \vee p^2)$.

$$p = \infty \dots$$

Approximate feature maps

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') \approx \Phi_{\mathbf{M}}(\mathbf{x})^\top \Phi_{\mathbf{M}}(\mathbf{x}')$$

$$\hat{K} \approx \underbrace{\hat{\Phi}_M^\top}_{M < n} \hat{\Phi}_M$$

Least squares with features

Let $\hat{\Phi}_M \in \mathbb{R}^{nM}$ with $(\hat{\Phi}_M)_{ij} = \Phi_M(\mathbf{x}_i)^j$

Then,

$$\hat{\mathbf{w}}_\lambda = (\hat{\Phi}_M^\top \hat{\Phi}_M + \lambda n \mathbf{I})^{-1} \hat{\Phi}_M^\top \hat{\mathbf{y}}$$

Requires: time $O(nM^2 + M^3)$ space $O(M^2)$.

Example: linear sketching

Let S be a $d \times M$ matrix and

$$\hat{\Phi}_M = \hat{X}S$$

Equivalently

$$\mathbf{x} \in \mathbb{R}^d \quad \mapsto \quad \Phi_M(\mathbf{x}) = (\mathbf{s}_j^\top \mathbf{x})_{j=1}^M \in \mathbb{R}^M$$

with $\mathbf{s}_1, \dots, \mathbf{s}_M$ columns of S .

Random linear sketching

Let $\mathbf{s}_j \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ iid, then

$$\mathbf{x}^\top \mathbf{x}' = \frac{1}{M} \mathbb{E}[\Phi_M(\mathbf{x})^\top \Phi_M(\mathbf{x}')].$$

Random linear sketching

Let $\mathbf{s}_j \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ iid, then

$$\mathbf{x}^\top \mathbf{x}' = \frac{1}{M} \mathbb{E}[\Phi_M(\mathbf{x})^\top \Phi_M(\mathbf{x}')].$$

Proof

$$\frac{1}{M} \mathbb{E}[\Phi_M(\mathbf{x})^\top \Phi_M(\mathbf{x}')] = \frac{1}{M} \mathbb{E}\left[\sum_{j=1}^M \mathbf{x}^\top \mathbf{s}_j \mathbf{s}_j^\top \mathbf{x}'\right] = \frac{1}{M} \sum_{j=1}^M \mathbf{x}^\top \underbrace{\mathbb{E}[\mathbf{s}_j \mathbf{s}_j^\top]}_{\text{Identity}} \mathbf{x}' = \mathbf{x}^\top \mathbf{x}'.$$

Random linear sketching

Let $\mathbf{s}_j \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ iid, then

$$\mathbf{x}^\top \mathbf{x}' = \frac{1}{M} \mathbb{E}[\Phi_M(\mathbf{x})^\top \Phi_M(\mathbf{x}')].$$

Proof

$$\frac{1}{M} \mathbb{E}[\Phi_M(\mathbf{x})^\top \Phi_M(\mathbf{x}')] = \frac{1}{M} \mathbb{E}\left[\sum_{j=1}^M \mathbf{x}^\top \mathbf{s}_j \mathbf{s}_j^\top \mathbf{x}'\right] = \frac{1}{M} \sum_{j=1}^M \mathbf{x}^\top \underbrace{\mathbb{E}[\mathbf{s}_j \mathbf{s}_j^\top]}_{\text{Identity}} \mathbf{x}' = \mathbf{x}^\top \mathbf{x}'.$$

Note:

- Related to Johnson-Linderstrauss Lemma...

Linear random features

$$\mathbf{x}^\top \mathbf{x}' = \mathbb{E}[(\mathbf{x}^\top \mathbf{s})(\mathbf{s}^\top \mathbf{x}')]$$

Linear random features

$$\mathbf{x}^\top \mathbf{x}' = \mathbb{E}[(\mathbf{x}^\top \mathbf{s})(\mathbf{s}^\top \mathbf{x}')]$$

\Downarrow

$$\mathbf{x}^\top \mathbf{x}' \approx \frac{1}{M} \sum_{j=1}^M \mathbf{x}^\top \mathbf{s}_j \mathbf{s}_j^\top \mathbf{x}'$$

Linear random features

$$\mathbf{x}^\top \mathbf{x}' = \mathbb{E}[(\mathbf{x}^\top \mathbf{s})(\mathbf{s}^\top \mathbf{x}')]$$

\Downarrow

$$\mathbf{x}^\top \mathbf{x}' \approx \frac{1}{M} \sum_{j=1}^M \mathbf{x}^\top \mathbf{s}_j \mathbf{s}_j^\top \mathbf{x}' = \frac{1}{M} \Phi_M(\mathbf{x})^\top \Phi_M(\mathbf{x}')$$

with $\Phi_M(\mathbf{x}) = (\mathbf{s}_j^\top \mathbf{x})_{j=1}^M \in \mathbb{R}^M$

Random features

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') = \mathbb{E}[\psi(\mathbf{x}, \mathbf{s})\psi(\mathbf{x}', \mathbf{s})]$$

Random features

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') = \mathbb{E}[\psi(\mathbf{x}, \mathbf{s})\psi(\mathbf{x}', \mathbf{s})]$$

\Downarrow

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') \approx \frac{1}{M} \sum_{j=1}^M \psi(\mathbf{x}, \mathbf{s}_j)\psi(\mathbf{x}', \mathbf{s}_j)$$

Random features

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') = \mathbb{E}[\psi(\mathbf{x}, \mathbf{s})\psi(\mathbf{x}', \mathbf{s})]$$

\Downarrow

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') \approx \frac{1}{M} \sum_{j=1}^M \psi(\mathbf{x}, \mathbf{s}_j)\psi(\mathbf{x}', \mathbf{s}_j) = \frac{1}{M} \Phi_M(\mathbf{x})^\top \Phi_M(\mathbf{x}')$$

with $\Phi_M(\mathbf{x}) = (\psi(\mathbf{x}, \mathbf{s}_1), \psi(\mathbf{x}, \mathbf{s}_2), \dots, \psi(\mathbf{x}, \mathbf{s}_M))$.

Random Fourier features

Let $\mathbf{X} = \mathbb{R}$, $\mathbf{s}_j \sim \mathcal{N}(0, 1)$ iid and

$$\psi(\mathbf{x}, \mathbf{s}_j) = \underbrace{e^{i\mathbf{s}_j\mathbf{x}}}_{\text{complex exp.}}.$$

Random Fourier features

Let $\mathbf{X} = \mathbb{R}$, $s_j \sim \mathcal{N}(0, 1)$ iid and

$$\psi(\mathbf{x}, s_j) = \underbrace{e^{is_j x}}_{\text{complex exp.}}.$$

For $k(\mathbf{x}, \mathbf{x}') = e^{-|\mathbf{x} - \mathbf{x}'|^2 \gamma}$, then

$$k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[\psi(\mathbf{x}, \mathbf{s})\psi(\mathbf{x}', \mathbf{s})]$$

Random Fourier features

Let $\mathbf{X} = \mathbb{R}$, $s_j \sim \mathcal{N}(0, 1)$ iid and

$$\psi(\mathbf{x}, s_j) = \underbrace{e^{is_j \mathbf{x}}}_{\text{complex exp.}}.$$

For $k(\mathbf{x}, \mathbf{x}') = e^{-|\mathbf{x} - \mathbf{x}'|^2 \gamma}$, then

$$k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[\psi(\mathbf{x}, \mathbf{s})\psi(\mathbf{x}', \mathbf{s})]$$

Proof: from basic properties of the Fourier transform

$$e^{-|\mathbf{x} - \mathbf{x}'|^2 \gamma} = \text{const.} \int d\mathbf{s} \underbrace{e^{i\mathbf{s}\mathbf{x}}}_{\text{Inv. transf.}} \underbrace{e^{-i\mathbf{s}\mathbf{x}'}}_{\text{Transl.}} \underbrace{e^{-\frac{\mathbf{s}^2}{\gamma}}}_{\text{Tranf. of Gaussian}}.$$

Random Fourier features (cont.)

- ▶ The above reasoning immediately extends to $X = \mathbb{R}^d$.
- ▶ Using symmetry one can show the same result holds for

$$\psi(\mathbf{x}, (\beta_j, \mathbf{b}_j)) = \cos(\beta_j^\top \mathbf{x} + \mathbf{b}_j)$$

with \mathbf{b}_j uniformly distributed.

Random ReLU features

Let $(\beta_j, \mathbf{b}_j) \sim \mathcal{U}[\mathbb{S}^{d+1}]$

$$\psi(\mathbf{x}, (\beta_j, \mathbf{b}_j)) = |\beta_j \mathbf{x} + \mathbf{b}_j|_+$$

If

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') = \sin \theta + (\pi - \theta) \cos \theta, \quad \theta = \arccos(\mathbf{x}^\top \mathbf{x}')$$

then

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') = \mathbb{E}[\psi(\mathbf{x}, \mathbf{s})\psi(\mathbf{x}', \mathbf{s})]$$

Kernels and random features

- Pick a kernel and derive a random feature map, or ...

Kernels and random features

- ▶ Pick a kernel and derive a random feature map, or ...
- ▶ ...pick random features and derive the limit kernel.

Kernels and random features

- ▶ Pick a kernel and derive a random feature map, or ...
- ▶ ...pick random features and derive the limit kernel.

Useful perspective for neural nets?

Random features and neural nets

$$f(\mathbf{x}) = \sum_{j=1}^M w_j \sigma(\beta_j^\top \mathbf{x} + \mathbf{b}_j)$$

- ▶ $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ a non linear activation function.
- ▶ For $j = 1, \dots, M$, β_j , w_j , \mathbf{b}_j parameters to be determined.

Random features and neural nets

$$f(\mathbf{x}) = \sum_{j=1}^M w_j \sigma(\beta_j^\top \mathbf{x} + b_j)$$

- ▶ $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ a non linear activation function.
- ▶ For $j = 1, \dots, M$, β_j , w_j , b_j parameters to be determined.

Some references

- ▶ **History** [McCulloch, Pitts '43; Rosenblatt '58; Minsky, Papert '69; Y. LeCun, '85; Hinton et al. '06]
- ▶ **Deep learning** [Krizhevsky et al. '12 - 18705 Cit.!!!]
- ▶ **Theory** [Barron '92-94; Bartlett, Anthony '99; Pinkus, '99]

Learning with random features

$$f(\mathbf{x}) = \sum_{j=1}^M w_j \sigma(\beta_j^\top \mathbf{x} + b_j)$$

- ▶ $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ a non linear activation (?) function.
- ▶ For $j = 1, \dots, M$, w_j parameters to be determined
- ▶ For $j = 1, \dots, M$, β_j, b_j **chosen at random**

Learning with random features

$$f(\mathbf{x}) = \sum_{j=1}^M w_j \sigma(\beta_j^\top \mathbf{x} + b_j)$$

- ▶ $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ a non linear activation (?) function.
- ▶ For $j = 1, \dots, M$, w_j parameters to be determined
- ▶ For $j = 1, \dots, M$, β_j, b_j **chosen at random**

Some references

- ▶ **Neural nets** [Block '62], **Extreme learning machine** [Huang et al. '06] 5196 Cit.??
- ▶ **Gaussian processes/kernel methods** [Neal '95, Rahimi, Recht '06'08'08, Acot et al. '18 (Neural tangent kernel)]

Mean field neural nets model

$$f(\mathbf{x}) = \sum_{j=1}^M w_j \sigma(\beta_j^\top \mathbf{x} + \mathbf{b}_j)$$

Mean field neural nets model

$$f(\mathbf{x}) = \sum_{j=1}^M w_j \sigma(\beta_j^\top \mathbf{x} + \mathbf{b}_j)$$

Infinitely wide neural nets define RKHS

$$f(\mathbf{x}) = \int d\pi(\beta, \mathbf{b}) w(\beta, \mathbf{b}) \sigma(\beta^\top \mathbf{x} + \mathbf{b})$$

Mean field neural nets model

$$f(\mathbf{x}) = \sum_{j=1}^M w_j \sigma(\beta_j^\top \mathbf{x} + \mathbf{b}_j)$$

Infinitely wide neural nets define RKHS

$$f(\mathbf{x}) = \int d\pi(\beta, \mathbf{b}) w(\beta, \mathbf{b}) \sigma(\beta^\top \mathbf{x} + \mathbf{b}) = \mathbb{E}[w(\beta, \mathbf{b}) \sigma(\beta^\top \mathbf{x} + \mathbf{b})].$$

...

Least squares and random features

$$\hat{K} \approx \underbrace{\hat{\Phi}_M}_{M < n} \hat{\Phi}_M^\top$$

$$\hat{\mathbf{w}}_\lambda = (\hat{\Phi}_M^\top \hat{\Phi}_M + \lambda n \mathbf{I})^{-1} \hat{\Phi}_M^\top \hat{\mathbf{y}} \quad \text{Requires: time } O(nM^2 + M^3) \quad \text{space } O(nM).$$

Even better: SGD

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \gamma_t \Phi_M(\mathbf{x}_t)(\mathbf{w}_t^\top \Phi_M(\mathbf{x}_t) - y_t) \quad \text{Requires: time } O(nMt) \quad \text{space } O(M).$$

How many random features?

$$\boxed{\hat{K}} \approx \underbrace{\boxed{\hat{\Phi}_M}}_{M < n} \boxed{\hat{\Phi}_M^\top}$$

kernel approximation

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') \approx \Phi_{\mathbf{M}}(\mathbf{x})^\top \Phi_{\mathbf{M}}(\mathbf{x}')$$

vs

learning

$$\mathbb{E}[\ell(\mathbf{y}, \hat{\mathbf{w}}_\lambda \Phi_{\mathbf{M}}(\mathbf{x}))]$$

The number of RF needed for learning can be much smaller than $n!$

But...

The number of RF needed for learning can be much smaller than $n!$

But... there's no time today!

What about data dependent approximations?

Outline

Large scale

Random features

Nystrom

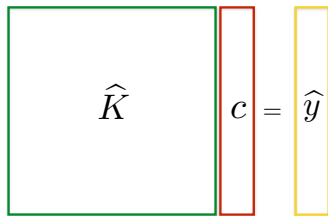
Optimization

Conclusion

Representer theorem

$$\hat{\mathbf{f}}_{\lambda}(\mathbf{x}) = \sum_{i=1}^n \mathbf{k}(\mathbf{x}, \mathbf{x}_i) \mathbf{c}_i$$

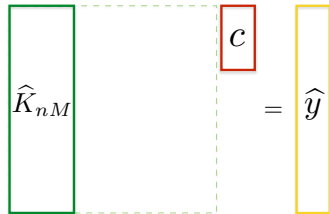
$$(\hat{K} + \lambda n \mathbf{I}) \mathbf{c} = \hat{\mathbf{y}}$$


$$\hat{K} \mathbf{c} = \hat{\mathbf{y}}$$

Nystrom

$$\hat{\mathbf{f}}_{\lambda, M}(\mathbf{x}) = \sum_{i=1}^M \mathbf{K}(\mathbf{x}, \tilde{\mathbf{x}}_i) \mathbf{c}_i$$

$$(\hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{K}}_{nM} + \lambda n \hat{\mathbf{K}}_{MM}) \mathbf{c} = \hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{y}}$$



$$\hat{K}_{nM} \mathbf{c} = \hat{\mathbf{y}}$$

with $\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_M \subset \mathbf{x}_1, \dots, \mathbf{x}_n$ and corresponding columns chosen at random

Nystrom

$$\hat{\mathbf{f}}_{\lambda, M}(\mathbf{x}) = \sum_{i=1}^M \mathbf{K}(\mathbf{x}, \tilde{\mathbf{x}}_i) \mathbf{c}_i$$

$$(\hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{K}}_{nM} + \lambda n \hat{\mathbf{K}}_{MM}) \mathbf{c} = \hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{y}}$$



with $\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_M \subset \mathbf{x}_1, \dots, \mathbf{x}_n$ and corresponding columns chosen at random

Not throwing data! Just using a subset for modeling functions!


Why should it work?

Why should it work?

$$\left(\underbrace{\widehat{\mathbf{X}}^\top \widehat{\mathbf{X}}}_{\text{Cov.matrix}} + \lambda \mathbf{n} \right)^{-1}$$

Why should it work?

$$\left(\underbrace{\hat{\mathbf{X}}^\top \hat{\mathbf{X}}}_{\text{Cov. matrix}} + \lambda \mathbf{n} \right)^{-1}$$



Why should it work?

$$\left(\underbrace{\hat{\mathbf{X}}^\top \hat{\mathbf{X}}}_{\text{Cov. matrix}} + \lambda \mathbf{n} \right)^{-1}$$



Why should it work?

$$\left(\underbrace{\hat{\mathbf{X}}^\top \hat{\mathbf{X}}}_{\text{Cov. matrix}} + \lambda \mathbf{n} \right)^{-1}$$



Why should it work?


$$(\hat{\mathbf{K}} + \lambda \mathbf{n})^{-1}$$



We apply the same idea to the feature/kernel space!

Name game

$$\hat{\mathbf{f}}_{\lambda, M}(\mathbf{x}) = \sum_{i=1}^M \mathbf{K}(\mathbf{x}, \tilde{\mathbf{x}}_i) \mathbf{c}_i$$
$$(\hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{K}}_{nM} + \lambda n \hat{\mathbf{K}}_{MM}) \mathbf{c} = \hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{y}}$$



Some references

- **Nystrom methods** (Smola, Scholkopf '00)
- Gaussian processes: inducing inputs (Quiñonero-Candela et al '05)
- Randomized numerical linear algebra: column sampling (Halko et al. '11)

Why the name Nyström approximation?

Why the name Nyström approximation?

Discrete approximation of integral operators

For all \mathbf{x}

$$\int \mathbf{k}(\mathbf{x}, \mathbf{x}') \mathbf{c}(\mathbf{x}') d\mathbf{x}' = \mathbf{y}(\mathbf{x}) \quad \mapsto \quad \sum_{j=1}^M \mathbf{k}(\mathbf{x}, \tilde{\mathbf{x}}_j) \mathbf{c}(\tilde{\mathbf{x}}_j) = \mathbf{y}(\tilde{\mathbf{x}}_j)$$

Related to **quadrature** methods.

Why the name Nyström approximation?

Discrete approximation of integral operators

For all \mathbf{x}

$$\int \mathbf{k}(\mathbf{x}, \mathbf{x}') \mathbf{c}(\mathbf{x}') d\mathbf{x}' = \mathbf{y}(\mathbf{x}) \quad \mapsto \quad \sum_{j=1}^M \mathbf{k}(\mathbf{x}, \tilde{\mathbf{x}}_j) \mathbf{c}(\tilde{\mathbf{x}}_j) = \mathbf{y}(\tilde{\mathbf{x}}_j)$$

Related to to **quadrature** methods.

From operators to (large)matrices

For all $i = 1, \dots, n$

$$\sum_{j=1}^n \mathbf{k}(\mathbf{x}_i, \mathbf{x}_j) \mathbf{c}_j = \mathbf{y}_i \quad \mapsto \quad \sum_{j=1}^M \mathbf{k}(\mathbf{x}_i, \tilde{\mathbf{x}}_j) \mathbf{c}_j = \mathbf{y}_i$$

Nyström approximation and subsampling

For all $i = 1, \dots, n$

$$\sum_{j=1}^n k(\mathbf{x}_i, \mathbf{x}_j) \mathbf{c}_j = \mathbf{y}_i \quad \mapsto \quad \sum_{j=1}^M k(\mathbf{x}_i, \tilde{\mathbf{x}}_j) \mathbf{c}_j = \mathbf{y}_i$$

The above formulation highlights connection to columns subsampling

$$\hat{\mathbf{K}} \mathbf{c} = \hat{\mathbf{y}} \quad \mapsto \quad \hat{\mathbf{K}}_{nM} \mathbf{c}_M = \hat{\mathbf{y}}$$

Nystrom as sketching

Consider the $d \times M$ matrix $S = (\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_M)$ and

$$\hat{\Phi}_M = \hat{X}S$$

Equivalently

$$\mathbf{x} \in \mathbb{R}^d \mapsto \Phi_M(\mathbf{x}) = (\tilde{\mathbf{x}}_j^\top \mathbf{x})_{j=1}^M \in \mathbb{R}^M$$

with $\mathbf{s}_1, \dots, \mathbf{s}_M$ columns of S .

Nystrom (for the linear kernels) is a form of data driven sketching.

Nystrom as projection regularization

The cost of Nyström kernel ridge regression

$$\hat{\mathbf{f}}_{\lambda, M}(\mathbf{x}) = \sum_{i=1}^M \mathbf{K}(\mathbf{x}, \tilde{\mathbf{x}}_i) \mathbf{c}_i$$

$$(\hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{K}}_{nM} + \lambda n \hat{\mathbf{K}}_{MM}) \mathbf{c} = \hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{y}}$$

$$\hat{K}_{nM} \mathbf{c} = \hat{\mathbf{y}}$$

Requires: time $O(nM^2 + M^3)$ space $O(nM)$.

How many Nyström centers?

$$\hat{K}_{nM} c = \hat{y}$$

kernel approximation

$$\hat{K} \approx \hat{K}_{nM} \hat{K}_{MM}^{-1} \hat{K}_{nM}^\top$$

vs

learning

$$\mathbb{E}[\ell(\mathbf{y}, \hat{\mathbf{f}}_{\lambda, M}(\mathbf{x}))]$$

The number of Nyström centers needed for learning can be much smaller than n !

But...

The number of Nyström centers needed for learning can be much smaller than n !

But... there's no time today!

Can we do better?

Possible improvements

- ▶ **adaptive sampling (leverage scores)**
- ▶ optimization

Leverage scores and sampling

$$\ell(\mathbf{i}, \lambda) = (\widehat{\mathbf{K}}(\widehat{\mathbf{K}} + \lambda \mathbf{nI})^{-1})_{ii}$$

Sampling $\mathbf{J} = \tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_{\mathbf{M}} \subset \mathbf{x}_1, \dots, \mathbf{x}_n$ according to $\ell(\mathbf{1}, \lambda), \dots, \ell(\mathbf{n}, \lambda)$.

- ▶ Can lead to smaller M than uniform sampling.
- ▶ Fast algorithms needed...

Approximate leverage scores

Basic idea: use only a subset of points to compute lev. scores.

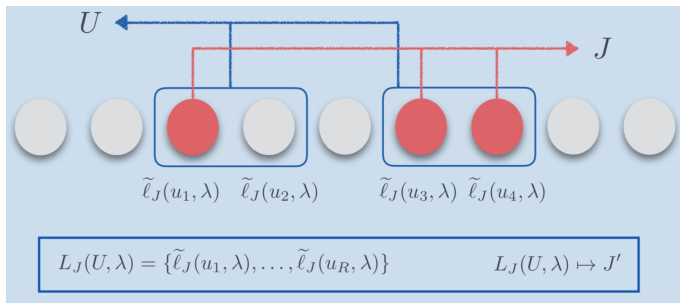
$$\tilde{\ell}(\mathbf{i}, \lambda) = \frac{1}{\lambda \mathbf{n}} (\hat{K}_{ii} - \hat{K}_{J,i}^\top (\hat{K}_{JJ} + \lambda \mathbf{n} \mathbf{W})^{-1} \hat{K}_{J,i})$$

Approximate leverage scores

Basic idea: use only a subset of points to compute lev. scores.


$$\tilde{\ell}(\mathbf{i}, \lambda) = \frac{1}{\lambda n} (\hat{\mathbf{K}}_{ii} - \hat{\mathbf{K}}_{J,i}^\top (\hat{\mathbf{K}}_{JJ} + \lambda n \mathbf{W})^{-1} \hat{\mathbf{K}}_{J,i})$$

Basic algorithm: uniform+lev. scores sampling.



Fast leverage scores sampling

BLESS algorithm: coarse to fine uniform+lev. scores sampling.



Possible improvements

- ▶ adaptive sampling (leverage scores)
- ▶ **optimization**

Outline

Large scale

Random features


Nystrom

Optimization

Conclusion

Beyond $O(n^2)$ time?


$$(\hat{K}_{nM}^\top \hat{K}_{nM} + \lambda n \hat{K}_{MM}) \mathbf{c} = \hat{K}_{nM}^\top \hat{\mathbf{y}}.$$



Bottleneck: computing $\hat{K}_{nM}^\top \hat{K}_{nM}$ requires $O(nM^2)$ time.

Optimization to rescue

$$\underbrace{\hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{K}}_{nM} + \lambda n \hat{\mathbf{K}}_{MM}}_{\mathbf{H}} \mathbf{c} = \underbrace{\hat{\mathbf{K}}_{nM}^\top \hat{\mathbf{y}}}_{\mathbf{b}}.$$



Idea: First order methods

$$\mathbf{c}_t = \mathbf{c}_{t-1} - \frac{\tau}{n} [\mathbf{K}_{nM}^\top (\mathbf{K}_{nM} \mathbf{c}_{t-1} - \mathbf{y}_n) + \lambda n \mathbf{K}_{MM} \mathbf{c}_{t-1}]$$

Pros: requires $O(nMt)$

Cons: $t \propto \kappa(\mathbf{H})$ arbitrarily large- $\kappa(\mathbf{H}) = \sigma_{\max}(\mathbf{H})/\sigma_{\min}(\mathbf{H})$ condition number.

Preconditioning

Idea: solve an equivalent linear system with better condition number

Preconditioning

$$Hc = b \quad \mapsto \quad P^T H P \beta = P^T b, \quad c = P \beta.$$

Ideally $PP^T = H^{-1}$, so that

$$t = O(\kappa(H)) \quad \mapsto \quad t = O(1)!$$

Computing a good preconditioning can be hard!

Remarks

- Preconditioning kernel ridge regression
(Fasshauer et al '12, Avron et al '16, Cutaját '16, Ma, Belkin '17)

$$H = \hat{K} + \lambda n I$$

Can we precondition Nystrom-KRR?

Preconditioning Nyström-KRR


Consider $H = \hat{K}_{nM}^\top \hat{K}_{nM} + \lambda n \hat{K}_{MM}$

Proposed Preconditioning

$$PP^\top = \left(\frac{n}{M} K_{MM}^2 + \lambda n K_{MM} \right)^{-1}$$

Compare to naive preconditioning

$$PP^\top = \left(\hat{K}_{nM}^\top \hat{K}_{nM} + \lambda n K_{MM} \right)^{-1}.$$



Baby FALKON

Proposed Preconditioning

$$\mathbf{P}\mathbf{P}^\top = \left(\frac{n}{M} \mathbf{K}_{\text{MM}}^2 + \lambda n \mathbf{K}_{\text{MM}} \right)^{-1},$$

Gradient descent

$$\hat{\mathbf{f}}_{\lambda, M, t}(\mathbf{x}) = \sum_{i=1}^M \mathbf{K}(\mathbf{x}, \tilde{\mathbf{x}}_i) \mathbf{c}_{t,i}, \quad \mathbf{c}_t = \mathbf{P} \beta_t$$

$$\beta_t = \beta_{t-1} - \frac{\tau}{n} \mathbf{P}^\top \left[\mathbf{K}_{nM}^\top (\mathbf{K}_{nM} \mathbf{P} \beta_{t-1} - \mathbf{y}_n) + \lambda n \mathbf{K}_{\text{MM}} \mathbf{P} \beta_{t-1} \right]$$

FALKON

- ▶ Gradient descent \mapsto conjugate gradient
- ▶ Computing P

$$P = \frac{1}{\sqrt{n}} T^{-1} A^{-1}, \quad T = \text{chol}(K_{MM}), \quad A = \text{chol}\left(\frac{1}{M} TT^T + \lambda I\right),$$

where $\text{chol}(\cdot)$ is the Cholesky decomposition.



Relevant works

References

- ▶ Less is more (Rudi et al. '15)
- ▶ Divide and conquer (Zhang et al. '13)
- ▶ NYTRO (Angles et al '16)
- ▶ Nyström SGD (Lin, R. '17)
- ▶ EIGENPRO (Belkin al '16)

Computational costs for FALKON




$\log n$ iterations suffice leading to

Space: $O(n)$
Kernel eval.: $O(nM)$
Time: $O(nM \log n)$
Model: $O(M)$

In practice

Higgs dataset: $n = 10,000,000$, $M = 50,000$



Some experiments

	MillionSongs (n ~ 10 ⁶)			YELP (n ~ 10 ⁶)		TIMIT (n ~ 10 ⁶)	
	MSE	Relative error	Time(s)	RMSE	Time(m)	c-err	Time(h)
FALKON	80.30	4.51×10^{-3}	55	0.833	20	32.3%	1.5
Prec. KRR	-	4.58×10^{-3}	289 [†]	-	-	-	-
Hierarchical	-	4.56×10^{-3}	293*	-	-	-	-
D&C	80.35	-	737*	-	-	-	-
Rand. Feat.	80.93	-	772*	-	-	-	-
Nyström	80.38	-	876*	-	-	-	-
ADMM R. F.	-	5.01×10^{-3}	958 [†]	-	-	-	-
BCD R. F.	-	-	-	0.949	42 [‡]	34.0%	1.7 [‡]
BCD Nyström	-	-	-	0.861	60 [‡]	33.7%	1.7 [‡]
KRR	-	4.55×10^{-3}	-	0.854	500 [‡]	33.5%	8.3 [‡]
EigenPro	-	-	-	-	-	32.6%	3.9 [‡]
Deep NN	-	-	-	-	-	32.4%	-
Sparse Kernels	-	-	-	-	-	30.9%	-
Ensemble	-	-	-	-	-	33.5%	-

Table: MillionSongs, YELP and TIMIT Datasets. Times obtained on: [‡] = cluster of 128 EC2 r3.2xlarge machines, [†] = cluster of 8 EC2 r3.8xlarge machines, [‡] = single machine with two Intel Xeon E5-2620, one Nvidia GTX Titan X GPU and 128GB of RAM, * = cluster with 512 GB of RAM and IBM POWER8 12-core processor, * = unknown platform.

Some more experiments

	SUSY ($n \sim 10^6$)			HIGGS ($n \sim 10^7$)		IMAGENET ($n \sim 10^6$)	
	c-err	AUC	Time(m)	AUC	Time(h)	c-err	Time(h)
FALKON	19.6%	0.877	4	0.833	3	20.7%	4
EigenPro	19.8%	-	6 [‡]	-	-	-	-
Hierarchical	20.1%	-	40 [†]	-	-	-	-
Boosted Decision Tree	-	0.863	-	0.810	-	-	-
Neural Network	-	0.875	-	0.816	-	-	-
Deep Neural Network	-	0.879	4680 [‡]	0.885	78 [‡]	-	-
Inception-V4	-	-	-	-	-	20.0%	-

Table: Architectures: [†] cluster with IBM POWER8 12-core cpu, 512 GB RAM, [‡] single machine with two Intel Xeon E5-2620, one Nvidia GTX Titan X GPU, 128GB RAM, [§] single machine.

Image classification


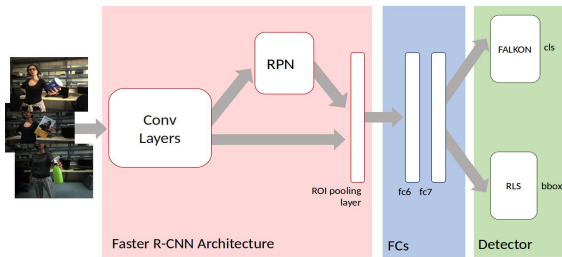
$$f(\mathbf{x}) = \langle \mathbf{w}, \Phi(\mathbf{x}) \rangle, \quad \mathbf{x} \mapsto \underbrace{\Phi_L}_{\text{Kernel representation}} \circ \underbrace{\Phi_{L-1} \cdots \circ \Phi_1(\mathbf{x})}_{\text{Convolutional}}$$

Imagenet

Top-1 class error	
FALKON + I-v4 feat.	20.7%
Inception-v4	20.0%
Inception-v3	21.2%
Inception-v2	23.4%
BN-Inception	25.2%
BN-GoogLeNet	26.8%
GoogLeNet	29.0%

Table: Single crop experimental results on the validation set of ILSVRC 2012.

Real time object detection in robotics



Method	mAP [%]	Train Time
Faster R-CNN	51,9	~25 min
FALKON + Full Bootstrap ($\sim 1K \times 1000$)	51,5	~8 min
FALKON + Random BKG (0×7000)	47,7	~25 sec



Take home message

You can train an SVM on millions of points in seconds/minutes

Outline

Large scale

Random features

Nystrom

Optimization

Conclusion

Who cares about kernels?